BoolNet package vignette

Christoph Müssel, Martin Hopfensitz, Hans A. Kestler

February 25, 2015

Contents

1 Introduction

`BoolNet` is an R package that provides tools for assembling, analyzing and visualizing synchronous and asynchronous Boolean networks as well as probabilistic Boolean networks. This document gives an introduction to the usage of the software and includes examples for use cases.

`BoolNet` supports four types of networks:

Synchronous Boolean networks consist of a set of Boolean variables

$$X = \{X_1, \ldots, X_n\}$$

and a set of transition functions

$$F = \{f_1, \ldots, f_n\},$$

one for each variable. These transition functions map an input of the Boolean variables in X to a Boolean value (0 or 1). We call a Boolean vector $\mathbf{x}(t) = (x_1(t), \ldots, x_n(t))$ the *state* of the network at time t. Then, the next state of the network $\mathbf{x}(t)$ is calculated by applying *all* transition functions $f_i(\mathbf{x}(t-1))$.

In a biological context, genes can be modeled as Boolean variables (*active/expressed* or *inactive/not expressed*), and the transition functions model the dependencies among these genes. In the synchronous model, the assumption is that all genes are updated at the same time. This simplification facilitates the analysis of the networks.

Asynchronous Boolean networks have the same structure as synchronous Boolean networks. Yet, at each point of time t, only *one* of the transition functions $f_i \in F$ is chosen at random, and the corresponding Boolean variable is updated. This corresponds to the assumption that in a genetic network, gene expression levels are likely to change at different points of time. In the most common model, the gene to be updated is chosen uniformly among all genes. Moreover, `BoolNet` supports specifying non-uniform update probabilities for the genes.

Probabilistic Boolean networks (PBN) allow for specifying more than one transition function per variable/gene. Each of these functions has a probability to be chosen, where the probabilities of all functions for one variable sum up to 1. Formally

$$F = \{\{(f_{11}, p_{11}), \ldots, (f_{1k_1}, p_{1k_1})\}, \ldots, \{(f_{n1}, p_{n1}), \ldots, (f_{nk_n}, p_{nk_n})\}\}$$

where k_i is the number of alternative transition functions for variable i, and p_{ij} is the probability that function j is chosen for variable i. A state transition is performed by selecting one function for each gene based on the probabilities and applying the chosen functions synchronously.

Temporal Boolean networks are Boolean networks that incorporate temporal predicates and discrete time delays. Here, the next state $x(t)$ may not only depend on $x(t-1)$, but can depend on any predecessor state $x(t-\Delta), \Delta \in \{1, 2, \ldots\}$. Furthermore, $x(t)$ may also directly depend on the time step t itself.

In `BoolNet`, there are different structure classes representing these network types:

***BooleanNetwork* objects** contain synchronous and asynchronous Boolean networks. Here, the transition functions are internally represented as truth tables.

***ProbabilisticBooleanNetwork* objects** encode Probabilistic Boolean networks. They use a truth table representation as well.

***SymbolicBooleanNetwork* objects** represent synchronous and temporal Boolean networks. They encode Boolean functions in a symbolic form, i.e. as expression trees.

As we have seen, the networks are represented in two different forms: The truth table representation, which basically maps inputs to the corresponding output values, is usually the most efficient representation for synchronous, asynchronous and probabilistic networks and uses a very fast simulator. However, this representation grows exponentially with the number of inputs and is therefore inappropriate for networks with a high number of inputs. This is particularly the case for temporal networks, where each unique time delay for a gene encodes an input. Hence, temporal networks are represented by directly encoding the corresponding Boolean expressions and use a different simulator. As synchronous Boolean networks are a special case of temporal networks (with all time delays being 1), these networks can also be represented as *SymbolicBooleanNetwork* objects.

The package provides several methods of constructing networks: Networks can be loaded from files in which human experts describe the dependencies between the genes. Furthermore, they can be reconstructed from time series of gene expression measurements. It is also possible to generate random networks. This can be helpful for the identification of distinct properties of biological networks by comparison to random structures. The different methods of assembling networks are described in Section 2.

In Section 3, tools for the analysis and visualization of network properties are introduced. For synchronous, asynchronous and temporal Boolean networks, the most important tool is the identification of attractors. Attractors are cycles of states and are assumed to be associated with the stable states of cell function. Another possibility of identifying relevant states is the included Markov chain simulation. This method is particularly suited for probabilistic networks and calculates the probability that a state is reached after a certain number of iterations. To test the robustness of structural properties of the networks to noise and mismeasurements, the software also includes extensive support for perturbing networks. In this way, it is possible to test these properties in noisy copies of a biological network.

In Section 4, the interaction of `BoolNet` with related software is described. In particular, the import from and export to SBML is discussed. Also, the necessary steps to import networks import networks from BioTapestry and to export networks to Pajek are outlined.

For the examples in the following sections, we assume that the `BoolNet` package has been properly installed into the R environment. This can be done by typing

```
> install.packages("BoolNet")
```

into the R console or by the corresponding menu entries in an R GUI. For some of the plots, the `igraph` package is required and must be installed in your R environment as well. This is analogous to installing `BoolNet`. For the BioTapestry and SBML import, the `XML` package must be installed. After installation, the `BoolNet` package can be loaded via

```
> library(BoolNet)
```

2 Assembling networks

2.1 Assembling a network from expert knowledge

A major advantage of Boolean networks is the fact that natural-language statements can easily be transferred into this representation. This allows researchers for building Boolean networks entirely from expert knowledge, for example by collecting statements on gene dependencies from literature and expressing them as Boolean rules.

`BoolNet` is able to read in networks consisting of such rule sets in a standardized text file format. In such a file, each line consists of a target gene and an update rule, usually separated by a comma.

Optionally, it is also possible to add a probability for the rule if the file describes a probabilistic network. The first line of such a file is a header

```
targets, factors
```

or

```
targets, factors, probabilities
```

To illustrate the process of transforming natural-language statements into Boolean rules, we take a look at the mammalian cell cycle network introduced by Fauré et al. [7]. In Table 1 of this paper, the authors list natural-language statements of gene dependencies and the corresponding Boolean expressions. The following rules are taken from this table.

For gene CycD, Fauré et al. state:

> *CycD is an input, considered as constant.*

Transforming this into a Boolean rule is rather simple: CycD does not change its value, which means that its value after a transition only depends on its previous value. Thus, the transition rule is

```
CycD, CycD
```

Gene Rb has a more complex description:

> *Rb is expressed in the absence of the cyclins, which inhibit it by phosphorylation [...]; it can be expressed in the presence of CycE or CycA if their inhibitory activity is blocked by p27.*

As a general rule, inhibition can be represented by a Boolean negation. In the `BoolNet` file format, a negation is expressed by the ! character. The referenced cyclins comprise the genes CycA, CycB, CycD, and CycE. If *all* these genes are absent, Rb is expressed – i.e. if CycA is not expressed *and* CycB is not expressed *and* CycD is not expressed *and* CycE is not expressed. A logical AND is embodied by the & character. Consequently, the first part of the rule is

```
 ! CycA & ! CycB & ! CycD & ! CycE
```

In combination with the above statement, the fact that Rb can be expressed in the presence of CycE and CycA if p27 is active means that CycB and CycD must not be active. Thus, the second part of the rule is

```
p27 & ! CycB & ! CycD
```

This statement is an exception (or alternative) to the first statement; this can be expressed as a logical OR, for which the | character is used.

The complete rule for gene Rb is thus

```
Rb, (! CycA & ! CycB & ! CycD & ! CycE) | (p27 & ! CycB & ! CycD)
```

After processing all genes in the table in this way, we get the following network description:

```
targets, factors
CycD, CycD
Rb, (! CycA & ! CycB & ! CycD & ! CycE) | (p27 & ! CycB & ! CycD)
E2F, (! Rb & ! CycA & ! CycB) | (p27 & ! Rb & ! CycB)
CycE, (E2F & ! Rb)
CycA, (E2F & ! Rb & ! Cdc20 & ! (Cdh1 & UbcH10)) | (CycA & ! Rb & ! Cdc20 & ! (Cdh1 & UbcH10))
p27, (! CycD & ! CycE & ! CycA & ! CycB) | (p27 & ! (CycE & CycA) & ! CycB &! CycD)
Cdc20, CycB
Cdh1,(! CycA & ! CycB) | (Cdc20) | (p27 & ! CycB)
UbcH10, ! Cdh1 | (Cdh1 & UbcH10 & (Cdc20 | CycA | CycB))
CycB, ! Cdc20 & ! Cdh1
```

Now save this description to a file "cellcycle.txt" in your R working directory. The network can be loaded via

```
> cellcycle <- loadNetwork("cellcycle.txt")
```

The result is an object of class *BooleanNetwork* containing a truth table representation of the network.

The same network is also included in `BoolNet` as an example and can be accessed via

```
> data(cellcycle)
```

`BoolNet` also provides several convenience operators that can be used to express complex Boolean functions in a compact way, e.g.

- `maj(a, b, ...)` is 1 if the majority of its operands are 1. Similarly, `sumgt(a, b, ..., N)` is 1 if more than N operands are 1, `sumlt(a, b, ..., N)` is 1 if less than N operands are 1, and `sumis(a, b, ..., N)` is 1 if exactly N operands are 1.

- `all(a, b, ...)` is 1 if all its operands are 1 (i.e. a logical AND), and `any(a, b, ...)` is 1 if at least one of its operands is 1 (i.e. a logical OR).

The cell cycle network is a classical Boolean network, where each transition function only depends on the previous state of the network. E.g., `CycB, ! Cdc20 & ! Cdh1` can be written formally as $CycB(t) = \neg Cdc20(t-1) \wedge \neg Cdh1(t-1)$. As already discussed before, `BoolNet` also incorporates several temporal extensions. For example, a transition function can also depend on the states of genes at earlier time points:

```
a, b[-3] & b[-2] & b
```

is 1 if b has been active in the previous three time steps. The operators described above can also incorporate time ranges. The previous statement can be written in a more compact way using the `all` operator:

```
a, all[d=-3..-1](b[d])
```

This defines a time delay variable `d` that can be used for time specifications inside the operator. It can also be used in arithmetic operations. E.g.,

```
a, all[d=-3..-1](b[d] & c[d-1])
```

specifies that a is active if b has been active in the previous three time steps and c has been active at time $t-4$, $t-3$ and $t-2$.

Apart from relative time specifications, the `BoolNet` network language also incorporates predicates that depend on the absolute time, i.e. the number of time steps that have elapsed since the initial state.

For example,

```
a, timeis(3)
```

specifies that `a` is active at time step 3 and inactive at all other time steps. Similarly, the predicates `timelt` and `timegt` evaluate to 1 before and after the specified time point respectively.

As the above examples do not cover all possibilities of the network description language, a full language specification is provided in Section 5.

For temporal networks, `BoolNet` uses a special symbolic simulator that represents the functions as expression trees, whereas the standard simulator is based on a truth table representation. These simulators are discussed in Section 3. As synchronous Boolean networks are a special case of temporal networks, they can also be simulated with the symbolic simulator. When a network is loaded from a file using `loadNetwork()`, the user can specify the parameter `symbolic=TRUE` to load it in form of a *SymbolicBooleanNetwork* object instead of a *BooleanNetwork* object. The same parameter is also available for the import functions discussed in Section 4. Temporal networks can only be loaded with `symbolic=TRUE`, as `BoolNet` cannot represent them as truth tables.

As many network generation and modification routines (such as random network generation and network reconstruction that are discussed in the following sections) internally use the truth table representation, there are conversion routines `truthTableToSymbolic()` and `symbolicToTruthTable()` that convert synchronous Boolean networks of class *BooleanNetwork* in a truth table representation to networks of class *SymbolicBooleanNetwork* in a symbolic representation and vice versa. For more details, please refer to the manual.

2.2 Reconstructing a network from time series

An entirely different approach of assembling a network is to infer rules from series of expression measurements of the involved genes over time. For example, microarray experiments can be conducted at different points of time to cover the expression levels of different cell states. To reconstruct networks from such data, `BoolNet` includes two reconstruction algorithms for synchronous Boolean networks, Best-Fit Extension [12] and REVEAL [14]. REVEAL requires the inferred functions to match the input time series perfectly, hence it is not always able to reconstruct networks in the presence of noisy and inconsistent measurements. Best-Fit Extension retrieves a set of functions with minimum error on the input and is thus suited for noisy data.

In the following, we introduce a tool chain for the reconstruction of a Probabilistic Boolean Network from time series using Best-Fit extension.

Microarray measurements are usually represented as matrices of real-valued numbers which, for example, quantify the expression levels of genes. `BoolNet` includes a real-valued time series of gene measurements from a project to analyze the yeast cell cycle [18] which can be loaded using

```
> data(yeastTimeSeries)
```

This data contains four preselected genes and a series of 14 measurements for each of these genes.

In a first step, the real-valued dataset has to be converted to binary data as required by the reconstruction algorithm. `BoolNet` offers several binarization algorithms in the function `binarizeTimeSeries()`. We here employ the default method which is based on k-means clustering (with $k = 2$ for active and inactive):

```
> binSeries <- binarizeTimeSeries(yeastTimeSeries)
```

The returned structure in `binSeries` has an element `$binarizedMeasurements` containing the binary time series, and, depending on the chosen binarization method, some other elements describing parameters of the binarization.

To reconstruct the network from this data, we call the Best-Fit Extension algorithm:

```
> net <- reconstructNetwork(binSeries$binarizedMeasurements,
+                           method="bestfit",
+                           maxK=4)
```

Here, `maxK` is the maximum number of input genes for a gene examined by the algorithm. The higher this number, the higher is the runtime and memory consumption of the reconstruction. We can now take a look at the network using

```
> net
```

```
Probabilistic Boolean network with 4 genes

Involved genes:
Fkh2 Swi5 Sic1 Clb1

Transition functions:

Alternative transition functions for gene Fkh2:
Fkh2 = <f(Clb1){01}> (error: 1)
Fkh2 = <f(Fkh2){01}> (error: 1)

Alternative transition functions for gene Swi5:
Swi5 = <f(Clb1){01}> (error: 1)
Swi5 = <f(Fkh2){01}> (error: 1)

Alternative transition functions for gene Sic1:
Sic1 = <f(Sic1,Clb1){0001}> (error: 1)
Sic1 = <f(Swi5,Sic1){0001}> (error: 1)
Sic1 = <f(Fkh2,Sic1){0001}> (error: 1)

Alternative transition functions for gene Clb1:
Clb1 = <f(Clb1){01}> (error: 1)
Clb1 = <f(Fkh2){01}> (error: 1)
```

The dependencies among the genes in the network can be visualized using the `plotNetworkWiring()` function. In this graph, each gene corresponds to a vertex, and the inputs of transition functions correspond to edges.

```
> plotNetworkWiring(net)
```

plots a graph similar to that at the top of Figure 1. To use this function, you must install the `igraph` package.

A network that involved the same genes was examined by Kim et al. [11]. When comparing the wiring graph of our reconstructed network with the reference network presented in Figure 2 of this paper, one observes a very high similarity between the two networks. However, the reconstructed network comprises too many links for the gene Sic1: The reference network does not contain a self-regulation of Sic1 or regulation of Sic1 by Fkh2. If it is known in advance that these regulations are not plausible, such prior knowledge can be supplied to the reconstruction algorithm:

```
> net <- reconstructNetwork(binSeries$binarizedMeasurements,
+                           method="bestfit",
+                           maxK=4,
+                           excludedDependencies = list("Sic1" = c("Sic1", "Fkh2")))
```

Figure 1: The wiring graph of the reconstructed network without prior knowledge (top) and with the inclusion of prior knowledge (bottom). Each node of the graph represents one gene, and each arrow represents a gene dependency.

The wiring of the reconstruction with prior knowledge is shown at the bottom of Figure 1. We can see that the two false links are now eliminated. Similar to `excludedDependencies`, there is also a parameter `requiredDependencies` that specifies dependencies that must be included in the network.

When `reconstructNetwork()` discovers multiple functions for a gene with the minimum error on the input data, it includes all of these functions as alternative functions with equal probability. Consequently, the function returns a `ProbabilisticBooleanNetwork` structure.

If you would like to obtain a `BooleanNetwork` object with only one function per gene from a probabilistic network, you can extract such a network by telling the software which of the functions you would like to use. This can be done by specifying the indices of the functions to extract:

```
> net <- reconstructNetwork(binSeries$binarizedMeasurements,
+ method="bestfit", maxK=4)
> singleNet <- chooseNetwork(net, c(1,2,3,2))
```

The result is a Boolean network that is created by extracting the first function of gene Fkh2, the second function of genes Swi5 and Clb1, and the third function of gene Sic1 from the above probabilistic network:

```
> singleNet

Boolean network with 4 genes

Involved genes:
Fkh2 Swi5 Sic1 Clb1

Transition functions:
Fkh2 = <f(Clb1){01}>
Swi5 = <f(Fkh2){01}>
Sic1 = <f(Fkh2,Sic1){0001}>
Clb1 = <f(Fkh2){01}>
```

`BoolNet` also supports the generation of artificial time series from existing networks: The `generateTimeSeries()` function generates a set of time series from a network using random start states and optionally adds Gaussian noise.

```
> series <- generateTimeSeries(cellcycle,
+                              numSeries=100,
+                              numMeasurements=10,
+                              noiseLevel=0.1)
```

generates a list of 100 time series by calculating 10 consecutive transitions from 100 randomly chosen network states in the mammalian cell cycle network. The series are subject to Gaussian noise with a standard deviation of 0.1, such that the result is a list of real-valued matrices.

We can now binarize these simulated measurements and try to reconstruct the original network:

```
> binSeries <- binarizeTimeSeries(series, method="kmeans")
> net <- reconstructNetwork(binSeries$binarizedMeasurements, method="bestfit")
> net

Probabilistic Boolean network with 10 genes

Involved genes:
CycD Rb E2F CycE CycA p27 Cdc20 Cdh1 UbcH10 CycB
```

```
Transition functions:

Alternative transition functions for gene CycD:
CycD = <f(CycD){01}> (error: 0)

Alternative transition functions for gene Rb:
Rb = <f(CycD,CycE,p27,CycB){1010001000000000}> (error: 0)

Alternative transition functions for gene E2F:
E2F = <f(Rb,CycA,p27,CycB){1010001000000000}> (error: 0)

Alternative transition functions for gene CycE:
CycE = <f(Rb,E2F){0100}> (error: 0)

Alternative transition functions for gene CycA:
CycA = <f(Rb,E2F,Cdc20,Cdh1,UbcH10){11000000111000000000000000000000}> (error: 5)
CycA = <f(Rb,E2F,CycA,Cdc20,UbcH10){00001100100011000000000*00000*00}> (error: 5)

Alternative transition functions for gene p27:
p27 = <f(CycD,CycE,CycA,p27,CycB){1010*010001000000000*00000000000}> (error: 0)

Alternative transition functions for gene Cdc20:
Cdc20 = <f(CycB){01}> (error: 0)

Alternative transition functions for gene Cdh1:
Cdh1 = <f(CycA,p27,Cdc20,CycB){1011101100111011}> (error: 0)

Alternative transition functions for gene UbcH10:
UbcH10 = <f(CycA,Cdc20,Cdh1,UbcH10,CycB){11110001111100111111001111*10011}> (error: 0)

Alternative transition functions for gene CycB:
CycB = <f(Cdc20,Cdh1){1000}> (error: 0)
```

Obviously, the number of generated time series is still to small to reconstruct the network unambiguously. However, the result comes close to the original network. We see that the functions of the network are not fully specified: At some positions, there are asterisks (∗) denoting a *don't care* value. This means that functions with a 0 and a 1 at this position match the time series equally well. That is, a partially defined function with m asterisks corresponds to 2^m fully defined Boolean functions. It is also possible to generate these fully defined functions instead of the partially defined function by setting the parameter `returnPBN` to true (which was the behaviour prior to `BoolNet` version 2.0). For many *don't care* values, this may consume a high amount of memory and computation time.

`generateTimeSeries()` can also generate time series with artificial knock-outs and overexpressions:

```
> series <- generateTimeSeries(cellcycle,
+                              numSeries=10,
+                              numMeasurements=10,
+                              perturbations=1,
+                              noiseLevel=0.1)
```

specifies that each generated time series is generated from a network where one randomly selected gene is artificially knocked down (constantly 0) or overexpressed (constantly 1). These perturbations are returned in an additional element `perturbations`.

```
> series$perturbations
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
CycD	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
Rb	NA	NA	NA	NA	NA	NA	0	NA	NA	NA
E2F	NA	0	NA	NA	NA	0	NA	NA	NA	NA
CycE	NA	NA	NA	NA	NA	NA	NA	NA	NA	1
CycA	NA	NA	NA	NA	NA	NA	NA	1	NA	NA
p27	NA	NA	NA	1	NA	NA	NA	NA	NA	NA
Cdc20	NA	NA	NA	NA	1	NA	NA	NA	NA	NA
Cdh1	NA	NA	NA	NA	NA	NA	NA	NA	1	NA
UbcH10	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
CycB	1	NA	0	NA	NA	NA	NA	NA	NA	NA

Here, each column corresponds to the perturbations applied in one series. A value of 1 denotes an overexpression, a value of 0 denotes a knock-out, and an N/A value means that no perturbation was applied to this gene.

The `reconstructNetwork()` function also supports the reconstruction from such perturbation experiments if it is known which genes were perturbed. First, we store the series and the perturbation matrix in separate variables and binarize the data as before:

```
> perturbations <- series$perturbations
> series$perturbations <- NULL
> binSeries <- binarizeTimeSeries(series, method="kmeans")
```

Now, we can reconstruct the network by specifying the `perturbations` parameter:

```
> net <- reconstructNetwork(binSeries$binarizedMeasurements,
+                           method="bestfit",
+                           perturbations=perturbations)
> net

Probabilistic Boolean network with 10 genes

Involved genes:
CycD Rb E2F CycE CycA p27 Cdc20 Cdh1 UbcH10 CycB

Transition functions:

Alternative transition functions for gene CycD:
CycD = <f(CycD){01}> (error: 0)

Alternative transition functions for gene Rb:
Rb = <f(CycD,CycE,Cdh1,UbcH10){0111000000000000}> (error: 0)
Rb = <f(CycD,CycE,Cdc20,UbcH10){1101000000*00000}> (error: 0)
Rb = <f(CycD,CycE,CycA,CycB){100*0*0000000*00}> (error: 0)

Alternative transition functions for gene E2F:
E2F = <f(Rb,CycA,p27,CycB){101*001*0*0*****}> (error: 0)

Alternative transition functions for gene CycE:
```

```
CycE = <f(Rb,E2F){0100}> (error: 0)
CycE = <f(CycD,E2F){0001}> (error: 0)


Alternative transition functions for gene CycA:
CycA = <f(Rb,E2F,CycA,Cdc20,UbcH10){0000110010**10**00*0****0*00****}> (error: 0)


Alternative transition functions for gene p27:
p27 = <f(CycD,CycE,Cdh1,UbcH10){0111000000000000}> (error: 0)
p27 = <f(CycD,CycE,Cdc20,UbcH10){1101000000*000*0}> (error: 0)
p27 = <f(CycD,CycE,CycA,CycB){100*0*0000000*00}> (error: 0)


Alternative transition functions for gene Cdc20:
Cdc20 = <f(CycB){01}> (error: 0)


Alternative transition functions for gene Cdh1:
Cdh1 = <f(CycA,p27,Cdc20,CycB){10111*1100111*1*}> (error: 0)
Cdh1 = <f(CycA,p27,Cdc20,Cdh1){01*1*1110011111*}> (error: 0)


Alternative transition functions for gene UbcH10:
UbcH10 = <f(CycA,Cdc20,Cdh1,UbcH10){1*001101110111*1}> (error: 0)


Alternative transition functions for gene CycB:
CycB = <f(Cdc20,Cdh1){1000}> (error: 0)
```

As we generated only 10 series in this case, the reconstructed network is much more incomplete than in the previous reconstruction.

In biological settings, perturbation experiments are probably one of the most frequent ways of exploring the behaviour of a regulatory network, as it is much easier to obtain various different responses by applying perturbations than by just measuring the wild type behaviour.

2.3   Creating random networks

To study structural properties of Boolean networks and to determine the specific properties of biological networks in comparison to arbitrary networks, it is often desirable to generate artificial networks. BoolNet comprises a facility for the generation of random N-K networks [9, 10]. In the standard N-K networks, N is the total number of genes, and K is the number of input genes for each gene transition function. Such a network can be generated using

```
> net <- generateRandomNKNetwork(n=10, k=3)
```

This creates a network with 10 genes, each of which has a transition function that depends on 3 genes and whose output is generated uniformly at random. Similarly, one can also specify different numbers of input genes for each gene:

```
> net <- generateRandomNKNetwork(n=10, k=c(1,2,3,1,3,2,3,2,1,1))
```

BoolNet does not only support this standard case, but allows for different methods of choosing the numbers of input genes (parameter topology), the input genes themselves (parameter linkage), and the transition functions (parameter functionGeneration). In the following, some examples are presented.

The command

```
> net <- generateRandomNKNetwork(n=20, k=20, topology="scale_free")
```

determines the numbers of input genes by drawing values from the scale-free Zeta distribution [1]. According to this distribution, most transition functions will have a small number of input genes, but a few transition functions may depend on a high number of genes. The shape of the Zeta distribution can be customized using an additional parameter `gamma`, which potentially increases the number of input genes when chosen small and vice versa.

```
> net <- generateRandomNKNetwork(n=10, k=3, linkage="lattice")
```

creates a network in which the transition functions of the genes depend on a choice of genes with adjacent indices [2]. This leads to networks with highly interdependent genes.

It is also possible to influence the truth tables of the functions in several ways. The parameter `zeroBias` changes the ratio of 1 and 0 returned by the functions:

```
> net <- generateRandomNKNetwork(n=10,
+                                k=3,
+                                functionGeneration="biased",
+                                zeroBias=0.75)
```

generates a network in which the outcome of a transition function is 0 for around 75% of the inputs.

A more intricate way of influencing the function generation is the specification of generation functions. Generation functions explicitly generate functions according to a specific function class. *Canalyzing functions* are assumed to occur frequently in biological systems [8]. A canalyzing function has the property that one input can determine the output value on its own, i.e. if this input is either active or inactive, the output of the function is always the same. *Nested canalyzing functions* are a recursive definition of canalyzing functions, where the part of the function that does not depend on the canalyzing input is a canalyzing function for another input. Such functions can be generated using the built-in generation functions `generateCanalyzing()` and `generateNestedCanalyzing()`:

```
> net1 <- generateRandomNKNetwork(n=10,
+                                 k=3,
+                                 functionGeneration=generateCanalyzing,
+                                 zeroBias=0.75)
> net2 <- generateRandomNKNetwork(n=10,
+                                 k=3,
+                                 functionGeneration=generateNestedCanalyzing,
+                                 zeroBias=0.75)
```

It is also possible to define own generation functions: A generation function receives a vector `input` of input gene indices as a parameter and returns a truth table result column with `2^length(input)` values representing the function.

If no explicit generation scheme is known for the function class of interest, validation functions can be used instead of generation functions. Validation functions verify whether randomly generated functions belong to a specific function class and reject invalid functions. Naturally, this is much less efficient than generating appropriate functions directly. An example is the generation of monotone functions, which are also thought to be biologically plausible. These function account for the assumption that a transcription factor usually either activates or inhibits a specific target

gene, but does not change the type of regulation depending on other factors [16]. We can specify a simple validation function that checks whether a Boolean function is monotone:

```
> isMonotone <- function(input, func)
+ {
+   for (i in seq_len(ncol(input)))
+   # check each input gene
+   {
+     groupResults <- split(func, input[,i])
+     if (any(groupResults[[1]] < groupResults[[2]]) &&
+         any(groupResults[[1]] > groupResults[[2]]))
+       # the function is not monotone
+       return(FALSE)
+   }
+   return(TRUE)
+ }
```

Here, `input` is a matrix containing the input part of the transition table, and `func` is the output of the Boolean function. In a monotone function, the values of the target may only change in one direction when switching one input. Hence, a function for which switching the value of an input gene sometimes switches the target from active to inactive, but also sometimes switches it from inactive to active is not monotone. This is validated by comparing the two groups of transition table entries for which the current input is active and inactive respectively.

If a validation function is supplied to `generateRandomNKNetwork()`, the generator generates Boolean functions until either the validation function returns `TRUE` or the maximum number of iterations (specified by the parameter `failureIterations`) is reached, in which case it fails.

```
> net <- generateRandomNKNetwork(n=10,
+                                 k=3,
+                                 validationFunction="isMonotone",
+                                 failureIterations=1000)
```

creates a network with 10 genes in which all functions have 3 inputs, of which at least one is canalyzing.

By default, `generateRandomNKNetwork()` creates functions that cannot be simplified, i.e. that do not contain any genes that are irrelevant for the outcome of the function. If desired, this behaviour can be changed by setting `noIrrelevantGenes` to `FALSE`.

The presented parameters can be combined, and there are further options and parameters, so that a broad variety of networks with different structural properties can be generated. For a full reference of the possible parameters, please refer to the manual.

2.4   Knock-out and overexpression of genes

`BoolNet` allows for temporarily knocking out and overexpressing genes in a network without touching the transition functions. This means that genes can be set to a fixed value, and in any calculation on the network, this fixed value is taken instead of the value of the corresponding transition function. Knocked-out and overexpressed genes speed up the analysis of the network, as they can be ignored in many calculations. For example, to knock out CycD in the mammalian cell cycle network, we call

```
> data(cellcycle)
> knockedOut <- fixGenes(cellcycle, "CycD", 0)
```

or alternatively use the gene index

```
> knockedOut <- fixGenes(cellcycle, 1, 0)
```

This sets the gene constantly to 0. To over-express the gene (i.e. to fix it to 1), the corresponding call is

```
> overExpressed <- fixGenes(cellcycle, "CycD", 1)
```

The command

```
> originalNet <- fixGenes(knockedOut, "CycD", -1)
```

reactivates the gene (for both knock and overexpression) and resets the network to its original state.

The function also accepts multiple genes in a single call, such as

```
> newNet <- fixGenes(cellcycle, c("CycD","CycE"), c(0,1))
```

which knocks out CycD and overexpresses CycE.

3 Network analysis

3.1 Simulation of state transitions

To simulate a state transition and identify successor states of a given state, BoolNet includes the function stateTransition(). The function supports transitions for all four types of networks.

The following code performs a synchronous state transition for the state in which all genes are set to 1 on the mammalian cell cycle network:

```
> data(cellcycle)
> stateTransition(cellcycle, rep(1,10))
```

CycD	Rb	E2F	CycE	CycA	p27	Cdc20	Cdh1	UbcH10	CycB
1	0	0	0	0	0	1	1	1	0

To calculate all state transitions in a synchronous network until an attractor is reached, you can call

```
> path <- getPathToAttractor(cellcycle, rep(0,10))
> path
```

	CycD	Rb	E2F	CycE	CycA	p27	Cdc20	Cdh1	UbcH10	CycB
1	0	0	0	0	0	0	0	0	0	0
2	0	1	1	0	0	1	0	1	1	1
3	0	0	0	0	0	0	1	0	1	0
4	0	1	1	0	0	1	0	1	1	0
5	0	1	0	0	0	1	0	1	0	0

The returned matrix consists of the subsequent states until an attractor is reached. Depending on the optional parameter includeAttractorStates, the sequence comprises all attractor states, only the first attractor state or none of the attractor states.

A sequence can be visualized by plotting a table of state changes:
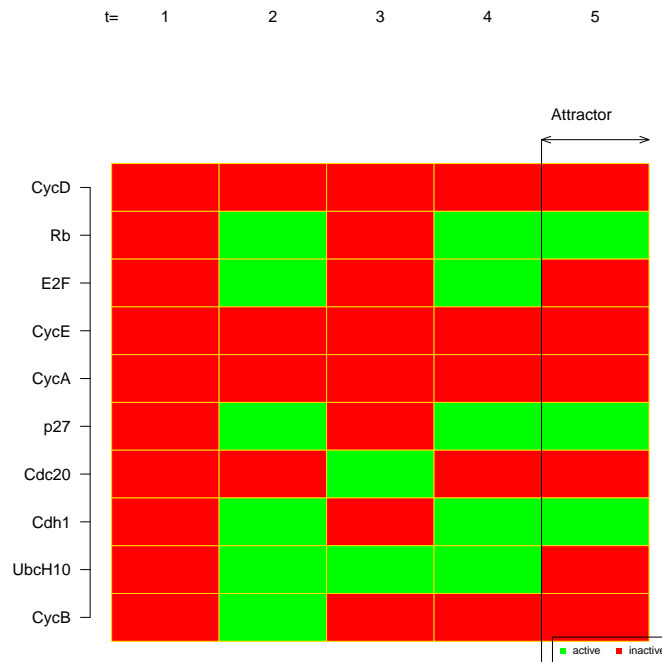
```
> plotSequence(sequence=path)
```

Figure 2: Visualization of a sequence of states in the mammalian cell cycle network. The columns of the table represent consecutive states of the time series. The last state is the steady-state attractor of the network.

The result is depicted in Figure 2. `plotSequence()` also includes a shortcut that calculates the sequence directly if a network and a start state are supplied. It also provides an alternative way of visualizing the sequence as a state transition by setting `mode="graph"`.

The function

```
> sequenceToLaTeX(sequence=path, file="sequence.tex")
```

creates a LaTeX table similar to `plotSequence()` function document.

In many cases, start states are defined by a set of active genes. Instead of supplying a full state vector, one can also supply only these active genes using the `generateState()` function.

```
> startState <- generateState(cellcycle, specs=c("CycD"=1,"CycA"=1))
> stateTransition(cellcycle,startState)
```

CycD	Rb	E2F	CycE	CycA	p27	Cdc20	Cdh1	UbcH10	CycB
1	0	0	0	1	0	0	0	1	1

calculates a state transition starting from a state where only the genes *CycD* and *CycA* are active, while all other genes are inactive (which is controlled by the `default` parameter of `generateState()`).

For temporal Boolean networks (objects of class *SymbolicBooleanNetworks*), the above functions can be utilized mostly in the same way. We demonstrate this using a small temporal network example of the IGF (Insulin-like growth receptor) pathway that is included in the package and can be loaded via

```
> data(igf)
```

The model illustrates the activation and feedback inhibition of the PI3K-Akt-mTOR signalling cascade through IGF and IRS.

16

A state transition from the initial state in which the trigger of the pathway – IGF – is active, can be performed using

```
> startState <- generateState(igf, specs=c("IGF"=1))
> stateTransition(igf, startState)

   IGF    IRS    PI3K    Akt mTORC1 mTORC2
     1      1      0      0      0      0
```

The IGF network incorporates time delays of up to 3. Therefore, the last three states have to be known to calculate a successor state. If only a single state is supplied – as above – the function assumes that the state was the same before. This can be seen when calculating the sequence to the attractor using

```
> getPathToAttractor(network=igf,state=startState)

         IGF IRS PI3K Akt mTORC1 mTORC2
t = -2    1   0    0   0      0      0
t = -1    1   0    0   0      0      0
t = 0     1   0    0   0      0      0
t = 1     1   1    0   0      0      0
t = 2     1   1    1   0      0      0
t = 3     1   1    1   1      0      1
t = 4     1   1    1   1      0      1
t = 5     1   1    1   1      0      0
t = 6     1   1    1   1      1      0
t = 7     1   1    1   1      1      0
t = 8     1   0    1   1      1      0
t = 9     1   0    0   1      1      0
t = 10    1   0    0   0      1      0
t = 11    1   0    0   0      1      0
t = 12    1   0    0   0      0      0
t = 13    1   0    0   0      0      0
t = 14    1   0    0   0      0      0
```

Here, the states at $t = -2, \ldots 0$ are the same as the generated start state. If it is required to specify multiple predecessor states here, a matrix of states can be supplied instead of a vector. E.g.,

```
> startState <- generateState(igf, specs=list("IGF"=c(0,0,1)))
> startState

      IGF IRS PI3K Akt mTORC1 mTORC2
[1,]    0   0    0   0      0      0
[2,]    0   0    0   0      0      0
[3,]    1   0    0   0      0      0
```

specifies that the first two states ($t = -2, t = -1$) should have all genes inactive, while IGF is activated only at $t = 0$. This time, we plot the sequence instead of printing it:

```
> plotSequence(network=igf, startState=startState)
```

The result is depicted in Figure 3.

stateTransition() can also perform asynchronous updates. A random asynchronous transition is performed using
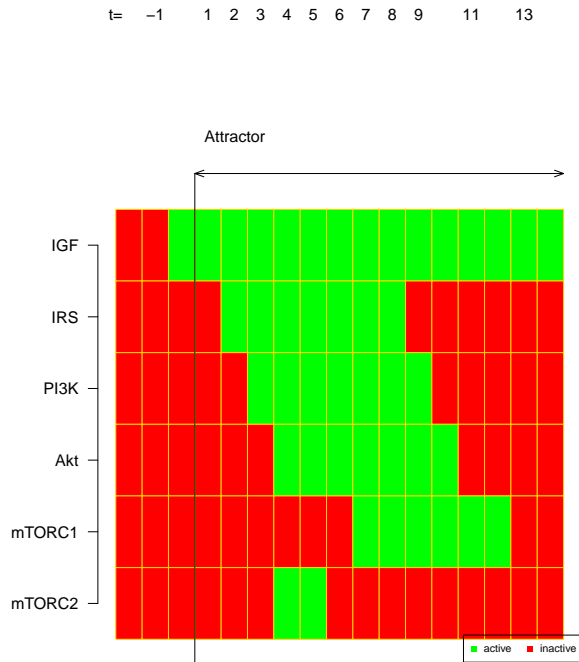
Figure 3: Visualization of a sequence of states in the IGF network. The columns of the table represent consecutive states of the time series. After activation of IGF, the attractor consisting of 14 states is entered immediately. This attractor represents the activation and inactivation of the PI3K-Akt-mTOR signalling cascade through IGF and IRS.

```
> stateTransition(cellcycle, rep(1,10), type="asynchronous")

  CycD    Rb   E2F  CycE  CycA   p27 Cdc20  Cdh1 UbcH10  CycB
     1     1     1     1     0     1     1     1      1     1
```

In this case, the fifth gene, CycA, was chosen at uniformly at random and updated.

We can also specify non-uniform probabilities for the genes, for example

```
> stateTransition(cellcycle, rep(1,10), type="asynchronous",
+ geneProbabilities=c(0.05,0.05,0.2,0.3,0.05,0.05,0.05,0.05,0.1,0.1))

  CycD    Rb   E2F  CycE  CycA   p27 Cdc20  Cdh1 UbcH10  CycB
     1     1     1     0     1     1     1     1      1     1
```

This obviously increases probabilities for the genes 3 and 4 (E2F and CycE) to be chosen. In this case, CycE was chosen for the update.

Sometimes you do not want a random update at all, but would like to specify which gene should be chosen for the update. This is possible via

```
> stateTransition(cellcycle, rep(1,10), type="asynchronous",
+ chosenGene="CycE")

  CycD    Rb   E2F  CycE  CycA   p27 Cdc20  Cdh1 UbcH10  CycB
     1     1     1     0     1     1     1     1      1     1
```

In probabilistic Boolean networks, a state transition is performed by choosing one of the alternative functions for each gene and applying this set of functions to the current state. The following performs a state transition with a randomly chosen set of functions on the artificial probabilistic Boolean network taken from [17] with 3 genes, starting from state (0,1,1):

18

```
> data(examplePBN)
> stateTransition(examplePBN, c(0,1,1), type="probabilistic")

x1 x2 x3
 1  0  0
```

You may get a different result, as the functions are chosen randomly according to the probabilities stored in the network. If you would like to execute a specific set of transition functions, you can supply this in an additional parameter:

```
> stateTransition(examplePBN, c(0,1,1), type="probabilistic",
+ chosenFunctions=c(2,1,2))

x1 x2 x3
 0  0  0
```

This call uses the second function for gene x1 and x3 and the first function for gene x2.

3.2 Identification of attractors

Attractors are stable cycles of states in a Boolean network. As they comprise the states in which the network resides most of the time, attractors in models of gene-regulatory networks are expected to be linked to phenotypes [10, 13]. Transitions from all states in a Boolean network eventually lead to an attractor, as the number of states in a network is finite. All states that lead to a certain attractor form its *basin of attraction*. BoolNet is able to identify attractors in synchronous and asynchronous Boolean networks. There are three types of attractors in these networks:

Simple attractors occur in synchronous and temporal Boolean networks and consist of a set of states whose synchronous transitions form a cycle.

Complex or loose attractors are the counterpart of simple attractors in asynchronous networks. As there is usually more than one possible transition for each state in an asynchronous network, a complex attractor is formed by two or more overlapping loops. Precisely, a complex attractor is a set of states in which all asynchronous state transitions lead to another state in the set, and a state in the set can be reached from all other states in the set.

Steady-state attractors are attractors that consist of only one state. All transitions from this state result in the state itself. These attractors are the same both for synchronous and asynchronous update of a network. Steady states are a special case of both simple attractors and complex attractors.

The getAttractors() function incorporates several methods for the identification of attractors in synchronous and asynchronous networks. We present these methods using the included mammalian cell cycle network as an example. This network has one steady-state attractor, one simple synchronous attractor consisting of 7 states, and one complex asynchronous attractor with 112 states (see [7]).

We first demonstrate the use of exhaustive synchronous state space search. This means that the software starts from all possible states of the network and performs synchronous state transitions until a simple or steady-state attractor is reached.

```
> data(cellcycle)
> attr <- getAttractors(cellcycle)
> attr

Attractor 1 is a simple attractor consisting of 1 state(s) and has a basin of 512
 state(s):

 |--<---------|
 V           |
 0100010100  |
 V           |
 |-->---------|


Genes are encoded in the following order: CycD Rb E2F CycE CycA p27 Cdc20 Cdh1
UbcH10 CycB

Attractor 2 is a simple attractor consisting of 7 state(s) and has a basin of 512
 state(s):

 |--<---------|
 V           |
 1001100000  |
 1000100011  |
 1000101011  |
 1000001110  |
 1010000110  |
 1011000100  |
 1011100100  |
 V           |
 |-->---------|


Genes are encoded in the following order: CycD Rb E2F CycE CycA p27 Cdc20 Cdh1
UbcH10 CycB
```

Typing `attr` calls a special print method that presents the attractor in a human-readable way. Here, a state in an attractor is represented by a binary vector, where each entry of the vector codes for one gene. An alternative is to print only the names of the active genes (i.e., the genes that are set to 1) instead of the full vector by calling the `print()` method explicitly with a changed parameter:

```
> print(attr, activeOnly=TRUE)

Attractor 1 is a simple attractor consisting of 1 state(s) and has a basin of 512
 state(s).
Active genes in the attractor state(s):
State 1: Rb, p27, Cdh1

Attractor 2 is a simple attractor consisting of 7 state(s) and has a basin of 512
 state(s).
Active genes in the attractor state(s):
State 1: CycD, CycE, CycA
State 2: CycD, CycA, UbcH10, CycB
State 3: CycD, CycA, Cdc20, UbcH10, CycB
State 4: CycD, Cdc20, Cdh1, UbcH10
State 5: CycD, E2F, Cdh1, UbcH10
State 6: CycD, E2F, CycE, Cdh1
State 7: CycD, E2F, CycE, CycA, Cdh1
```

Attractors with 7 state(s)

Figure 4: Visualization of the state changes in an attractor. The columns of the table represent consecutive states of the attractor. On top, the percentage of states leading to the attractor is supplied.

We can see that the search identified both synchronous attractors.

The `AttractorInfo` structure stores the attractors in an encoded form. The function `getAttractorSequence()` can be used to obtain the sequence of states that constitute a specific synchronous attractor as a table:

```
> getAttractorSequence(attr, 2)
```

```
  CycD Rb E2F CycE CycA p27 Cdc20 Cdh1 UbcH10 CycB
1    1  0   0    1    1   0     0    0      0    0
2    1  0   0    0    1   0     0    0      1    1
3    1  0   0    0    1   0     1    0      1    1
4    1  0   0    0    0   0     1    1      1    0
5    1  0   1    0    0   0     0    1      1    0
6    1  0   1    1    0   0     0    1      0    0
7    1  0   1    1    1   0     0    1      0    0
```

retrieves the states that make up the second (i.e., the 7-state attractor) as a data frame with the genes in the columns and the successive states in the rows.

Synchronous attractors can also be visualized by plotting a table of changes of gene values in the states of the attractor:

```
> plotAttractors(attr, subset=2)
```

plots the state changes of the simple attractor with 7 states, as depicted in Figure 4. Similarly,

```
> attractorsToLaTeX(attr, subset=2, file="attractors.tex")
```

21

exports the same state table to a LATEX document.

The advantage of the exhaustive search method is that apart from the attractors themselves, the complete transition table is calculated and stored in the return value. This table comprises information that is used by a number of analysis methods described below.

You can extract the transition table in a data frame and print it out using

```
> tt <- getTransitionTable(attr)
> tt

      State       Next state  Attr. basin  # trans. to attr.
0000000000 =>   0110010111            1                    4
[...]
1111111111 =>   1000001110            2                    1

Genes are encoded in the following order: CycD Rb E2F CycE
CycA p27 Cdc20 Cdh1 UbcH10 CycB
```

In the printed table, the first column denotes the initial state, the second column contains the state after the transition, the first column contains the number of the attractor that is finally reached from this state, and the fourth column lists the number of state transitions required to attain this attractor.

A table of the same structure is returned by

```
> getBasinOfAttraction(attr, 1)
```

which extracts all states from the transition table that belong to the basin of attraction of attractor one (i.e., whose attractor number in column 3 is 1).

If you are interested in information on a single state (here: the state with all genes set to 1), you can type

```
> getStateSummary(attr, c(1,1,1,1,1,1,1,1,1,1))

      State       Next state  Attr. basin  # trans. to attr.
1111111111 =>   1000001110            2                    1

Genes are encoded in the following order: CycD Rb E2F CycE
CycA p27 Cdc20 Cdh1 UbcH10 CycB
```

The visualization function `getStateGraph()` makes use of the transition table as well: It plots a transition graph in which the basins of attraction are drawn in different colors, and the attractors are highlighted. The result of

```
> plotStateGraph(attr)
```

is depicted at the top of Figure 5. The blue basin belongs to attractor 1, and the green basin belongs to attractor 2.

The above call does not ensure that the basins of attraction are clearly separated in the plot. If this is desired, one can choose to use a piecewise layout, which means that the layouting function is applied separately to each basin of attraction, and the basins are drawn side by side. The result of

```
> plotStateGraph(attr, piecewise=TRUE)
```
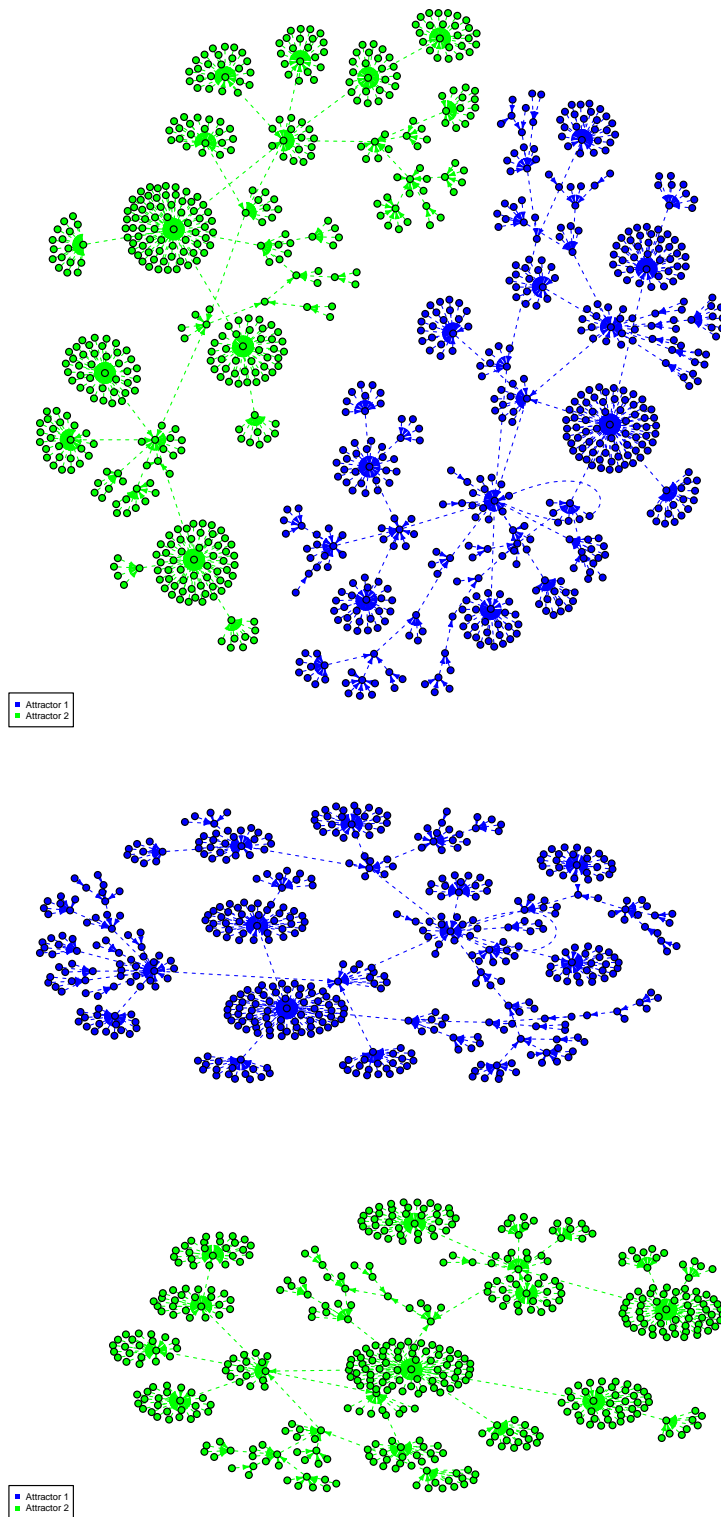
is depicted at the bottom of Figure 5.

Figure 5: The state graph of the mammalian cell cycle network using the regular layout (top) and using a piecewise layout (bottom). Each node represents a state of the network, and each arrow is a state transition. The colors mark different basins of attraction. Attractors are highlighted using bold lines.

Exhaustive search consumes a high amount of time and memory with increasing size of the network, which makes it intractable for large networks (`BoolNet` currently supports networks with up to 29 genes for exhaustive state space search). Therefore, `BoolNet` also allows for other search techniques that can handle larger networks. Depending on the requirements, the following options are available:

Heuristic state space search basically uses the same technique as exhaustive attractor search, but only uses a subset of states as initial states instead of the full state space. It then identifies the attractors to which state transitions from these states lead. The start states can either be supplied, or they can be calculated randomly.

Exhaustive SAT-based search exhaustively identifies all possible attractors of a network based on a formulation of attractor search as a satisfiability (SAT) problem and then applies the PicoSAT solver [4] to the problem. The algorithm is an adaption of that of Dubrova and Teslenko [6]. Unlike exhaustive state space search, this method only returns the attractors, but not the state transition table/graph. Depending on the network structure, it can be applied to networks comprising several hundreds of genes.

Restricted SAT-based search uses another formulation as a satisfiability problem and only identifies attractors of a predefined maximum cycle length. It is exhaustive with regard to attractors of at most this size, but will never return any attractors that comprise more than the predefined number of states. Like the exhaustive SAT-based search, this method does not return a transition table.

The following shows the application of the above methods exemplarily.

A heuristic state space search from randomly chosen start states in the mammalian cell cycle network can be started via

```
> attr <- getAttractors(cellcycle, method="random", startStates=100)
```

This chooses 100 random start states for the heuristic search and usually identifies both attractors.

```
> attr <- getAttractors(cellcycle,
+                       method="chosen",
+                       startStates=list(rep(0,10),rep(1,10)))
```

starts from the states (0,0,0,0,0,0,0,0,0,0) and (1,1,1,1,1,1,1,1,1,1) and again identifies both synchronous attractors.

For the above calls, only the subset of the transition table traversed by the heuristic search is returned. This means that there is no guarantee that, e.g. `getBasinOfAttraction()` returns the complete basin of attraction of an attractor in heuristic mode.

An exhaustive SAT-based search can be started using

```
> attr <- getAttractors(cellcycle,
+                       method="sat.exhaustive")
```

This will again return all attractors, but no transition table.

If you are only interested in a subset of attractors that have a predefined length, the restricted SAT-based search can be applied. For example,

```
> attr <- getAttractors(cellcycle,
+                       method="sat.restricted",
+                       maxAttractorLength=1)
```

restricts the attractor search to steady states. Correspondingly, only the first attractor of the network is returned:

```
> attr
```

Attractor 1 is a simple attractor consisting of 1 state(s):

```
 |--<---------|
 V           |
 0100010100  |
 V           |
 |-->---------|
```

Genes are encoded in the following order: CycD Rb E2F CycE CycA p27 Cdc20 Cdh1
UbcH10 CycB

BoolNet can also identify attractors in asynchronous Boolean network. For this purpose, another
special heuristic algorithm is included. This algorithm starts from a small subset of states and
makes a number of random transitions to reach an attractor with a high probability. After that,
a validation step is performed to analyze whether a complex attractor has been identified.

The command

```
> attr <- getAttractors(cellcycle,
+                       type="asynchronous",
+                       method="random",
+                       startStates=500)
```

conducts an asynchronous search with 500 random start states on the mammalian cell cycle net-
work. In this case, the algorithm has identified both the steady-state attractor and the complex
attractor:

```
> attr
```

Attractor 1 is a simple attractor consisting of 1 state(s):

```
 |--<---------|
 V           |
 0100010100  |
 V           |
 |-->---------|
```

Genes are encoded in the following order: CycD Rb E2F CycE
CycA p27 Cdc20 Cdh1 UbcH10 CycB

Attractor 2 is a complex/loose attractor consisting of 112 state(s)
and 338 transition(s):

```
1011101111 => 1011101110
[...]
1000000000 => 1010000000
```

Genes are encoded in the following order: CycD Rb E2F CycE
CycA p27 Cdc20 Cdh1 UbcH10 CycB

For the complex attractor, the involved transitions are printed out. By default, the algorithm
tries to avoid self-loops, i.e. transitions that lead to the same state again. This means that self-
loop transitions are only allowed if there is no other transition that leads to a different state. If
you would like to allow the algorithm to enter self-loops even if transitions to different states are
possible, you can call

Figure 6: Graph representation of the complex attractor in the mammalian cell cycle network. Each node represents a state of the complex attractor, and each arrow represents a state transition.

```
> attr <- getAttractors(cellcycle,
+                       type="asynchronous",
+                       method="random",
+                       startStates=500,
+                       avoidSelfLoops=FALSE)
```

In the resulting complex attractor with 112 states, there are 450 transitions instead of 338 transitions, which is due to the additional self-loops.

The asynchronous heuristic search does not return a transition table, such that the above analysis methods cannot be applied here.

As there are multiple possible transitions for each state, complex attractors cannot be visualized as in Figure 4. For this reason, plotAttractors() supports a graph mode that visualizes the transitions among the states in the attractor:

```
> plotAttractors(attr, subset=2, mode="graph", drawLabels=FALSE)
```

plots the 112-state attractor as depicted in Figure 6. We omit the state labels (i.e. the gene values) due to the high number of states. This plot again requires the igraph package.

Although getAttractors() can also be applied to temporal networks and other networks that are in a symbolic representation (i.e. *SymbolicBooleanNetwork* objects), this function is only a shortcut to the simulation function simulateSymbolicModel() in this case. It is advised to use simulateSymbolicModel() directly , as it provides more options. For the temporal model of the IGF pathway included in BoolNet, an exhaustive simulation can be performed as follows:

```
> sim <- simulateSymbolicModel(igf)
> sim


Simulation of a symbolic Boolean network
Sequences for 2048 start states (print with sequences=TRUE to show them)
```

```
Graph containing 320 state transitions (print with graph=TRUE to show them)

2 Attractors:
Attractor 1 is a simple attractor consisting of 1 state(s) and has a basin of 160
 state(s):

 |--<-----|
 V        |
 000000   |
 V        |
 |-->-----|


Genes are encoded in the following order: IGF IRS PI3K Akt mTORC1 mTORC2

Attractor 2 is a simple attractor consisting of 14 state(s) and has a basin of
160 state(s):

 |--<-----|
 V        |
 100000   |
 110000   |
 111000   |
 111101   |
 111101   |
 111100   |
 111110   |
 111110   |
 101110   |
 100110   |
 100010   |
 100010   |
 100000   |
 100000   |
 V        |
 |-->-----|


Genes are encoded in the following order: IGF IRS PI3K Akt mTORC1 mTORC2
```

By default, the result object of class *SymbolicSimulation* comprises several components:

- A list of sequences `sequences` from each start state to the corresponding attractor. If this component is not desired, the parameter `returnSequences` can be set to false.

- A graph structure `graph` that comprises all traversed state transitions. If this component is not desired, the parameter `returnGraph` can be set to false.

- The identified attractors `attractors`. If this component is not desired, the parameter `returnAttractors` can be set to false.

In this case, the network has two attractors: A steady state describes the inactive state of the pathway. The circular attractor describes the activation and inactivation of the PI3K-Akt-mTOR signalling cascade initiated by IGF.

All visualization and analysis function described above can also be applied to the simulation results obtained by `simulateSymbolicModel`. For example, the cascade attractor can be visualized via

```
> plotAttractors(sim, subset=2)
```

Similarly,

```
> plotStateGraph(sim)
```

plots the state transition graph of the network. Unlike in classical synchronous networks, a state can have multiple successor states (outgoing edges) in temporal networks, as a state transition may also depend on the history of states before the current state. The two plots are shown in Figure 7.

For temporal networks, it is often infeasible to perform an exhaustive state-space simulation, as the search space is not only exponential in the number of genes, but also in the time delays. There are two options to deal with this:

First, the search can be restricted to randomly generated or prespecified start states similarly to `getAttractors()`. If time delays of more than 1 are included in the network, not only single start states are generated, but the required history of states is generated as well. For example,

```
> sim <- simulateSymbolicModel(igf, method="random", startStates=2)
```

generates two start state matrices, each comprising 3 states (the maximum delay of the IGF network), and uses them as the basis for a simulation. This can be seen when examining the sequences to the attractors:

```
> sim$sequences
```

```
[[1]]
        IGF IRS PI3K Akt mTORC1 mTORC2
t = -2   1   1    1   0      0      0
t = -1   1   1    1   1      1      0
t = 0    0   1    1   1      0      0
t = 1    0   0    1   1      0      0
t = 2    0   0    0   1      1      0
t = 3    0   0    0   0      1      0
t = 4    0   0    0   0      1      0
t = 5    0   0    0   0      0      0
t = 6    0   0    0   0      0      0
t = 7    0   0    0   0      0      0


[[2]]
        IGF IRS PI3K Akt mTORC1 mTORC2
t = -2   1   0    0   0      1      0
t = -1   1   0    0   1      1      0
t = 0    0   0    0   1      1      0
t = 1    0   0    0   0      0      0
t = 2    0   0    0   0      1      0
t = 3    0   0    0   0      0      0
t = 4    0   0    0   0      0      0
t = 5    0   0    0   0      0      0
```

Both sequences comprise start states $t = -2$, $t = -1$ and $t = 0$.

The second option is to use a SAT-based attractor search similar to that in `getAttractors()`. Here, only the attractors are identified, but no state sequences and no state transition graph are returned.

```
> sim <- simulateSymbolicModel(igf, method="sat.exhaustive")
```

identifies all attractors in the network in an exhaustive manner. Alternatively, one can only search for attractors of a specific maximum length:
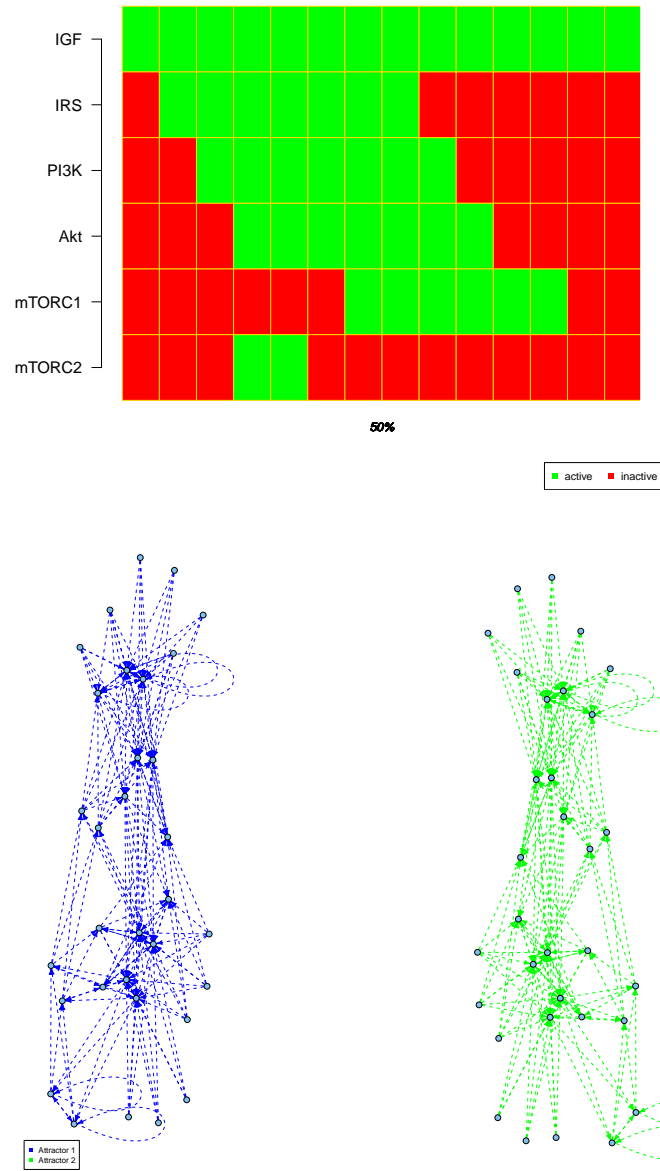
Figure 7: Top: Visualization of an attractor that describes the activation and inactivation of the PI3K-Akt-mTOR signalling cascade through IGF and IRS. The columns of the table represent consecutive states of the attractor. On top, the percentage of states leading to the attractor is supplied.

Bottom: The state transition graph of the IGF pathway network. Each node represents a state of the network, and each arrow is a state transition. The colors mark different basins of attraction.

```
> sim <- simulateSymbolicModel(igf, method="sat.restricted", maxAttractorLength=1)
```

will only return the steady-state attractor representing pathway inactivity, but not the 14-state attractor. Please note that this length-restricted search cannot be applied to temporal networks comprising the time-dependent predicates `timeis`,`timelt` or `timegt`.

Classical synchronous Boolean networks can also be simulated using the symbolic simulator `simulateSymbolicModel()` if they are in a symbolic form. However, in most cases, `getAttractors()` will be faster and will consume less memory for synchronous networks without temporal elements. Only if the number of inputs to genes is very high and exhaustive simulation is not required, it may be advisable to use the symbolic simulator.

3.3 Markov chain simulations

Another way of identifying relevant states in Boolean networks are Markov chain simulations. Instead of identifying cycles explicitly, these simulations calculate the probability that a certain state is reached after a predefined number of iterations. Of course, states in an attractor have a high probability of being reached if the number of iterations is chosen large enough. Markov chain simulations for probabilistic Boolean networks were introduced by Shmulevich et al. [17]. As a special case of probabilistic Boolean networks, these simulations are also suited for synchronous Boolean networks.

The following performs a Markov experiment with the predefined number of 1000 iterations on the example PBN described in [17]:

```
> data(examplePBN)
> sim <- markovSimulation(examplePBN)
> sim

States reached at the end of the simulation:
  x1 x2 x3 Probability
1  0  0  0        0.15
2  1  1  1        0.85


Probabilities of state transitions in the network:
  State      Next state  Probability
    000 =>          000          1.0
    001 =>          110          1.0
    010 =>          110          1.0
    011 =>          000          0.2
    011 =>          100          0.3
    011 =>          001          0.2
    011 =>          101          0.3
    100 =>          010          1.0
    101 =>          110          0.5
    101 =>          111          0.5
    110 =>          100          0.5
    110 =>          101          0.5
    111 =>          111          1.0
```

Only states with a non-zero probability are listed in the two tables. The first table shows the states that are reached after 1000 iterations. The second table is a transition table annotated with transition probabilities. This table can be suppressed by the parameter `returnTable=FALSE`. The results correspond exactly to those in [17].

If the transition table is included in the simulation results, we can plot a graph of the network:

```
> plotPBNTransitions(sim)
```

This graph is displayed in Figure 8. The vertices are the states of the graph. The edges represent transitions and are annotated with the corresponding transition probabilities. For this plot, the `igraph` package must be installed.
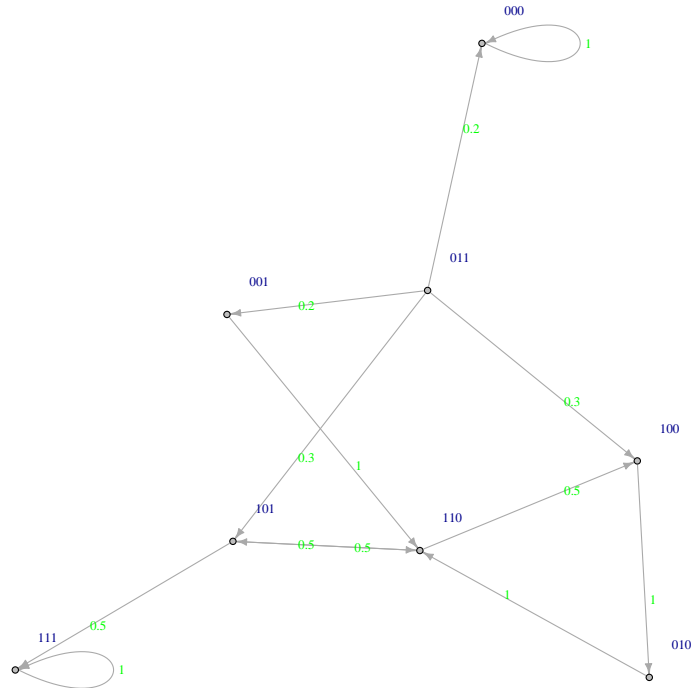

Figure 8: State transition graph of the example probabilistic Boolean network included in `BoolNet`. Each node represents a state of the network, and each arrow is a possible state transition, annotated by the transition probability.

We can also use Markov chain simulations to identify the attractor states in the mammalian cell cycle network:

```
> data(cellcycle)
> sim <- markovSimulation(cellcycle,
+                         numIterations=1024,
+                         returnTable=FALSE)
> sim
```

States reached at the end of the simulation:

	CycD	Rb	E2F	CycE	CycA	p27	Cdc20	Cdh1	UbcH10	CycB	Probability
1	1	0	0	1	1	0	0	0	0	0	0.00781250
2	1	0	1	1	0	0	0	1	0	0	0.17187500
3	1	0	1	1	1	0	0	1	0	0	0.02343750
4	0	1	0	0	0	1	0	1	0	0	0.50000000
5	1	0	1	0	0	0	0	1	1	0	0.15625000
6	1	0	0	0	0	0	1	1	1	0	0.10937500
7	1	0	0	0	1	0	0	0	1	1	0.00390625
8	1	0	0	0	1	0	1	0	1	1	0.02734375

We set the maximum number of iterations to 1024, which is the number of states in the network. In a deterministic network, this guarantees that all states are found.

The fourth state in the returned table is the steady-state attractor identified previously. It has a probability of 0.5, as the basin of attraction is exactly half of the states. The other 7 states belong to the simple synchronous attractor.

It is also possible to restrict the simulation to a certain set of input states instead of using all possible input states. In the following example, we only consider the state with all genes set to 1, and identify the state belonging to the steady-state attractor again:

```
> sim <- markovSimulation(cellcycle,
+                         numIterations=1024,
+                         returnTable=FALSE,
+                         startStates=list(rep(1,10)))
> sim
```

```
States reached at the end of the simulation:
  CycD Rb E2F CycE CycA p27 Cdc20 Cdh1 UbcH10 CycB Probability
1    1  0   1    0    0   0     0    1      1    0           1
```

3.4 Robustness assessment

A biological network is assumed to be robust to small amounts of noise. The plausibility of network models is therefore often assessed by testing its robustness to noise and mismeasurements. Typically, artificial noise is applied, and its influence on the behaviour of a network is measured.

There are two major ways of applying random noise: Either the current state of a network in a simulation can be perturbed, or the network structure itself can be perturbed. BoolNet includes functions for both types of robustness assessment.

The function perturbTrajectories measures the influence of noise that is applied to the current network state. It generates a set of initial states and creates perturbed copies of these states by randomly flipping bits. It then measures the influence of the flips on the further dynamic behaviour of the network. For example

```
> data(cellcycle)
> r <- perturbTrajectories(cellcycle,
+                          measure="hamming",
+                          numSamples=100,
+                          flipBits=1)
```

randomly generates 100 states and 100 copies with one bitflip and performs a single state transition for each state. It then measures the normalized Hamming distance (the fraction of different bits) between each state and the corresponding perturbed copy. A robust network is assumed to yield a low Hamming distance. The average distance can be viewed by typing

```
> r$value
```

```
[1] 0.107
```

A related measure is the average sensitivity. This measure assesses only a single transition function and counts the number of successor states that differ between the original states and the perturbed copies for the corresponding gene. E.g.,

```
> r <- perturbTrajectories(cellcycle,
+                          measure="sensitivity",
+                          numSamples=100,
+                          flipBits=1,
+                          gene="CycE")
```

```
> r$value
```

```
[1] 0.11
```

measures the average sensitivity of the transition function for gene CycE.

The long-term behaviour can be evaluated by comparing the attractors that are reached from the initial states and their perturbed copies.

```
> r <- perturbTrajectories(cellcycle,
+                          measure="attractor",
+                          numSamples=100,
+                          flipBits=1)
> r$value
```

```
[1] 0.93
```

measures the fraction of pairs of states and perturbed copies that yield the same attractors. It can be assumed that small changes in the state should not influence the long-term behaviour of the network, and hence the attractors should mostly be the same.

The second class of perturbations adds random noise to the network itself. This is implemented in the `perturbNetwork()` function. Unlike `perturbTrajectories()`, this function does not perform any simulations, but returns a perturbed copy of the network that can be analyzed further.

`BoolNet` includes a set of different perturbation options that can be combined. For example,

```
> perturbedNet <- perturbNetwork(cellcycle,
+                                perturb="functions",
+                                method="bitflip")
```

chooses a function of the network at random and flips a single bit in this function. By setting the parameter `maxNumBits`, you can also flip more than one bit at a time.

Instead of flipping bits,

```
> perturbedNet <- perturbNetwork(cellcycle,
+                                perturb="functions",
+                                method="shuffle")
```

randomly permutes the output values of the chosen transition functions. This preserves the numbers of 0s and 1s, but may change the Boolean function completely. These kinds of perturbations are supported for synchronous and asynchronous networks as well as for probabilistic networks.

For synchronous networks, a further perturbation mode is available:

```
> perturbedNet <- perturbNetwork(cellcycle,
+                                perturb="transitions",
+                                method="bitflip",
+                                numStates=10)
```

Here, `BoolNet` calculates the complete transition table of the network and then flips a single bit in 10 states of the transition table. From this modified table, a network is reconstructed. Changes of this type only affect a few states (which might not be the case when perturbing the functions directly as above), but possibly several of the transition functions. As in the previous examples, it is also possible to modify the number of bits to be flipped or to choose `method="shuffle"`.

A detailed listing of a perturbation experiment is shown below. In this experiment, 1000 perturbed copies of the cell cycle network are created, and the occurrences of the original synchronous attractors are counted in the perturbed copies. This is similar to the simulation in `perturbTrajectories()` with `measure="attractor"`. However, instead of comparing the outcomes of different initial states in the same network, it compares all attractors of different perturbed networks to all attractors of the original network by exhaustive search.

```
> # Perform a robustness test on a network
> # by counting the numbers of perturbed networks
> # containing the attractors of the original net
>
> library(BoolNet)
> # load mammalian cell cycle network
> data(cellcycle)
> # get attractors in original network
> attrs <- getAttractors(cellcycle, canonical=TRUE)
> # create 1000 perturbed copies of the network and search for attractors
> perturbationResults <- sapply(1:1000, function(i)
+ {
+   # perturb network and identify attractors
+   perturbedNet <- perturbNetwork(cellcycle, perturb="functions", method="bitflip")
+   perturbedAttrs <- getAttractors(perturbedNet, canonical=TRUE)
+
+   # check whether the attractors in the original network exist in the perturbed network
+   attractorIndices <- sapply(attrs$attractors,function(attractor1)
+       {
+          index <- which(sapply(perturbedAttrs$attractors, function(attractor2)
+             {
+                identical(attractor1, attractor2)
+             }))
+          if (length(index) == 0)
+            NA
+          else
+            index
+       })
+   return(attractorIndices)
+ })
> # perturbationResults now contains a matrix
> # with the first 2 columns specifying the indices or the
> # original attractors in the perturbed network
> # (or NA if the attractor was not found) and the next 2
> # columns counting the numbers of states
> # in the basin of attraction (or NA if the attractor was not found)
>
> # measure the total numbers of occurrences of the original attractors in the perturbed copies
> numOccurrences <- apply(perturbationResults[seq_along(attrs$attractors),,drop=FALSE], 1,
+                    function(row)sum(!is.na(row)))
> # print original attractors
> cat("Attractors in original network:\n")
> print(attrs)
> # print information
> cat("Number of occurrences of the original attractors",
+   "in 1000 perturbed copies of the network:\n")
> for (i in 1:length(attrs$attractors))
+ {
+   cat("Attractor ",i,": ",numOccurrences[i],"\n",sep="")
+ }
```

The results of such an experiment could look like this:

```
Attractors in original network:
Attractor 1 is a simple attractor consisting of 1 state(s)
and has a basin of 512 state(s):

 [...]


Attractor 2 is a simple attractor consisting of 7 state(s)
and has a basin of 512 state(s):

[...]


Number of occurrences of the original attractors in 1000
perturbed copies of the network:
Attractor 1: 622
Attractor 2: 589
```

We see that the steady-state attractor is slightly more robust to perturbations than the simple attractor with 7 states, as it can be identified in a higher number of perturbed copies.

3.5 Identifying specific properties of biological networks

The described robustness measures could also be used to identify specific properties of real-world networks in comparison to arbitrary (random) networks. For example, one could assume that attractors in biological networks are more robust to perturbations than attractors in random networks with a similar structure, as they should be capable of compensating for small dysfunctions of their components. Similarly to the above code, one could execute a number of random perturbations on the biological network and measure the percentage of original attractors found in the perturbed copies. Afterwards, one could repeat this process on a number of randomly generated networks – i.e., generate perturbed copies from each of the random networks, and measure the percentage of attractors in the copies. If the percentage of the biological network is higher than most of the percentages of the random network, this suggests that the biological network exhibits a higher robustness. This is a kind of computer-intensive test.

`BoolNet` comprises a generic facility for such computer-intensive tests. This facility already includes several tests (mainly for synchronous Boolean networks) and can be extended by custom test functions. The outlined example of attractor robustness is one of the integrated functions:

```
> data(cellcycle)
> res <- testNetworkProperties(cellcycle,
+                    numRandomNets=100,
+                    testFunction="testAttractorRobustness",
+                    testFunctionParams = list(copies=100, perturb="functions"))
```

creates a set of 100 random networks (each with the same number of input genes for the functions as the cell cycle network) and creates 100 perturbed copies for each of these networks and for the cell cycle network by applying `perturbNetwork()` with `perturb="functions"`. For each network, it then calculates the percentage of attractors that can still be found in the perturbed copies.

The function plots a histogram of this robustness measures of the random networks (see Figure 9, top panel). The corresponding value of the original cell cycle network is plotted as a red line, and the 95% quantile is plotted as a blue line.

We can see that the average percentage of found attractors is significantly higher in the biological network with a p-value of 0.01.

Figure 9: Top: Attractor robustness of randomly generated networks (histogram) in comparison to the mammalian cell cycle network (red line).
Bottom: Normalized Hamming distance of randomly generated networks (histogram) in comparison to the mammalian cell cycle network (red line).

It is also possible to perturb the states instead of the networks themselves by setting `perturb` to `"trajectories"`. In this case, the function applies `perturbTrajectories()` with `measure="attractor"` to the biological network and the randomly generated networks. It then tests whether the fraction of state pairs that yield the same attractor is higher in the biological network than in the randomly generated networks.

The second built-in test function also tests the robustness of the network behaviour by perturbing the network states: `testTransitionRobustness()` applies `perturbTrajectories()` with `measure="hamming"` to each network. It then checks whether random bit flips yield a higher Hamming distance of the successor states in randomly generated networks than in the biological model, i.e. whether noise in the states influences the randomly generated networks stronger than the biological model. In contrast to the previous measures for which a greater value was assumed in the biological model, the Hamming distance is assumed to be smaller. Hence, we must specify the test alternative as `alternative="less"`.

```
> testNetworkProperties(cellcycle,
+                       numRandomNets=100,
+                       testFunction="testTransitionRobustness",
+                       testFunctionParams=list(numSamples=100),
+                       alternative="less")
```

The results are shown in the bottom panel of Figure 9. Again, the result is highly significant (indeed, the Hamming distances are *always* lower in the biological network), which means that the biological network is considerably more robust to noise in the states than the randomly generated models.

Another network property can also be tested using a built-in function: When looking at the state graph of a biological network (which can be generated using `plotStateGraph()`), it can often be observed that many state transitions lead to the same successor states, which means that the dynamics of the network quickly concentrate on a few states after a number of state transitions. We call the number of states whose synchronous state transitions lead to a state s the *in-degree* of state s. We expect the biological network to have a few states with a high in-degree and many states with a low in-degree. A characteristic to summarize the in-degrees is the Gini index, which is a measure of inhomogeneity. If all states have an in-degree of 1, the Gini index is 0; if all state transitions lead to only one state, the Gini index is 1.

```
> testNetworkProperties(cellcycle,
+                       numRandomNets=100,
+                       testFunction="testIndegree")
```

plots an histogram of Gini indices in 100 random networks and draws the Gini index of the cell cycle network as a red line, as depicted in the top panel of Figure 10.

The histogram shows that the Gini index of the in-degrees is always higher in the biological network. This is probably due to the special structure of functions in biological networks.

Instead of accumulating the in-degrees using the Gini index, it is also possible to compare the distributions of the in-degrees across the networks. For this purpose, the Kullback-Leibler distances of the in-degrees of the supplied network and each of the random networks are calculated and plotted in a histogram. The Kullback-Leibler distance (also called relative entropy) is an asymetric measure of similarity of two distributions [5]. If the distributions are equal, the Kullback-Leibler distance is 0, otherwise it is greater than 0.

```
> testNetworkProperties(cellcycle,
+                       numRandomNets=100,
+                       testFunction="testIndegree",
+                       accumulation="kullback_leibler")
```

Figure 10: Top: Gini indices of state in-degrees of randomly generated networks (histogram) in comparison to the mammalian cell cycle network (red line)
Bottom: Kullback-Leibler distances of in-degrees of the mammalian cell cycle network and 100 random networks.

results in the plot displayed in the bottom panel of Figure 10.

It is possible to switch between the histogram of an accumulated characteristic (e.g. the Gini index) and the histogram of the Kullback-Leibler distances for all tests.

You can also easily implement your own tests. To do this, the only thing you have to do is implement a custom testing function that replaces `testIndegree()` or `testAttractorRobustness()`. Testing functions have the following signature:

```
function(network, accumulate=TRUE, params)
```

The first parameter is the network that should be tested. The parameter `accumulate` specifies whether a single characteristic value (e.g., the Gini index of the in-degrees) should be calculated, or whether a distribution of values (e.g., a vector of all in-degrees) should be returned. The third parameter is a list of further arguments needed by your function.

If, for example, we would like to compare the sizes of the basins of attractions of synchronous attractors in biological and random networks, we would write a function like this:

```
> testBasinSizes <- function(network, accumulate=TRUE, params)
+ {
+   attr <- getAttractors(network)
+   basinSizes <- sapply(attr$attractors, function(a)
+                       {
+                           a$basinSize
+                       })
+   if (accumulate)
+     return(mean(basinSizes))
+   else
+     return(basinSizes)
+ }
```

This function calculates the mean basin size as a characteristic value if accumulation is required, or returns the sizes of all basins of attraction in a vector otherwise. It does not need any further parameters in `params`.

Now, we can start a test using

```
> testNetworkProperties(cellcycle,
+                       numRandomNets=100,
+                       testFunction="testBasinSizes",
+                       xlab="Average size of basins of attraction")
```

to produce the plot shown in Figure 11. Apparently, the average basin sizes do not differ as much as the built-in test characteristics between the random networks and the cell cycle network.

By writing custom test functions, you can extend the test facility to perform a wide variety computer-intensive test. Of course, it is also possible to plot the Kullback-Leibler distances with such new methods by using `accumulation="kullback_leibler"`.

`testNetworkProperties()` accepts most of the parameters of `generateRandomNKNetwork()`. If necessary, you can generate more specialized kinds of random networks which resemble the original network in certain aspects, for example by specifying a generation function or by setting a proportion of 0 and 1 in the function outputs similar to the original network using `functionGeneration="biased"`.

Figure 11: A custom test statistic measuring the basin sizes on randomly generated networks (histogram) and the mammalian cell cycle network (red line).

4 Import and export

4.1 Saving networks in the `BoolNet` file format

Corresponding to the `loadNetwork()` command, a network can be saved using `saveNetwork()`. This stores the network in the network file format described in Section 5 and can be applied to all types of networks supported by `BoolNet`. For example, the cell cycle network can be saved using

```
> saveNetwork(cellcycle, file="cellcycle.txt")
```

The function stores the expressions that describe the transition functions. In some cases, there may not always be a valid symbolic description of the networks (e.g. for networks returned by `generateRandomNKNetwork()` when the `readableFunctions` parameter was not set). In this case, `saveNetwork()` can generate symbolic representations of the transition functions in Disjunctive Normal Form (DNF):

```
> net <- generateRandomNKNetwork(n=10, k=3, readableFunctions=FALSE)
> saveNetwork(net, file="randomnet.txt", generateDNF=TRUE)
```

The `generateDNF` parameter can also be used to detail which type of DNF formulae should be exported: `generateDNF="canonical"` exports canonical DNF formulae. `generateDNF="short"` minimizes the canonical functions by joining terms. By simply setting `generateDNF=TRUE`, formulae with up to 12 inputs are minimized, and formulae with more than 12 inputs are exported in a canonical form, as a minization is very time-consuming in this case.

4.2 Import from and export to SBML

`BoolNet` provides an interface to the widely used Systems Biology Markup Language (SBML) via the import function `loadSBML()` and the export function `saveSBML()`. As the core SBML does not fully support Boolean models, import and export of SBML models is based on the `sbml-qual` package which extends SBML by several qualitative modeling approaches, such as general logical models and Petri nets. For a full description of `sbml-qual`, refer to http://sbml.org/Documents/Specifications/SBML_Level_3/Packages/Qualitative_Models_(qual).

`BoolNet` only supports a subset of `sbml-qual`. It can read and write logical models with two possible values for each state, which are equivalent to Boolean networks. Logical models with more than two values for a gene or Petri nets cannot currently be handled by `BoolNet`.

An export to SBML is usually not associated with any loss of information. For example, we can write the cell cycle network to a file and re-import it into `BoolNet`:

```
> toSBML(cellcycle, file="cellcycle.sbml")
> sbml_cellcycle <- loadSBML("cellcycle.sbml")
> sbml_cellcycle

Boolean network with 10 genes

Involved genes:
CycD Rb E2F CycE CycA p27 Cdc20 Cdh1 UbcH10 CycB

Transition functions:
CycD = CycD
Rb = ((!CycA & !CycB & !CycD & !CycE) | (p27 & !CycB & !CycD))
E2F = ((!Rb & !CycA & !CycB) | (p27 & !Rb & !CycB))
CycE = (E2F & !Rb)
CycA = ((E2F & !Rb & !Cdc20 & !(Cdh1 & UbcH10)) | (CycA & !Rb & !Cdc20 & !(Cdh1 & UbcH10)))
p27 = ((!CycD & !CycE & !CycA & !CycB) | (p27 & !(CycE & CycA) & !CycB & !CycD))
Cdc20 = CycB
Cdh1 = ((!CycA & !CycB) | Cdc20 | (p27 & !CycB))
UbcH10 = (!Cdh1 | (Cdh1 & UbcH10 & (Cdc20 | CycA | CycB)))
CycB = (!Cdc20 & !Cdh1)
```

Apart from some additional brackets, the re-imported network coincides with the original network.

Similar to the `saveNetwork()` function, `toSBML()` exports a symbolic representation of the network transition functionss, which may not always be available. As for `saveNetwork()`, there is a parameter `generateDNF` that can be set to generate a symbolic representation in Disjunctive Normal Form from the truth tables.

4.3   Importing networks from BioTapestry

BioTapestry is a widely-used application for visual modeling of gene-regulatory networks [15]. It can be freely accessed at `http://www.biotapestry.org`. Although its primary purpose is visualization, the software supports specifying logical functions for the genes. `BoolNet` can read in the top-level ("Full genome") plot of a BioTapestry file (*.btp) and convert it into a Boolean network.

As an example, we assume the following BioTapestry model with 5 genes (2 inputs and 3 dependent genes):


The corresponding BioTapestry file is included in `BoolNet`. You can determine its path using

```
> system.file("doc/example.btp", package="BoolNet")
```

to access it in BioTapestry or `BoolNet`.

For the import, `BoolNet` needs to know the type of influence a gene has on another gene. Therefore, imported networks should only use links that are either enhancers or repressors. Neutral links are ignored in the import.

We now set further simulation parameters for the model. These parameters are imported by `BoolNet` to construct the functions of the Boolean network. First, we want to change the function of Gene 2 to `OR`. Right-click on Gene 2 and choose `Simulation Properties....`
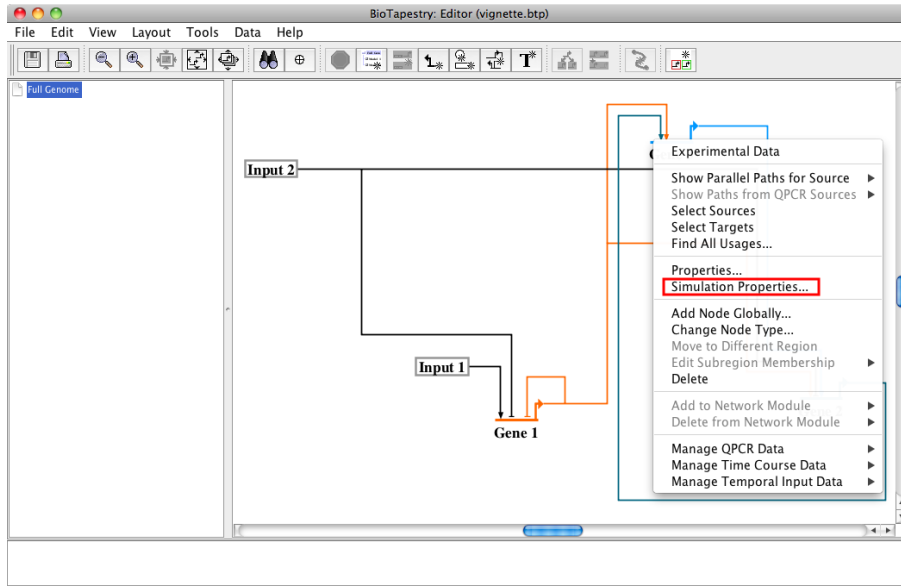
In the properties dialog, choose the `Logic` tab, and select `OR` for the logical function.

Now set the function of Gene 1 to `XOR` (exclusive or) in the same way.

You can also specify initial values for constant genes, i.e., genes with no input links. Choose the simulation properties of Input 1, and change to the `Parameters` tab. Choose `initVal` and set it to 1.

Press Return to store the result, and exit the dialog with `OK`. This will create a fixed gene with value 1 (i.e., an over-expressed gene) in the `BoolNet` import. Note that values other than 0 and 1 are ignored by the import, as well as initialization values for non-constant genes.

We assume that you save the network to a file "example.btp" in your working directory. In `R`, type

```
> net <- loadBioTapestry("example.btp")
```

to import the network. Alternatively, replace the file name by the command on page 42 to use the file in the package if you do not want to create the file yourself.

The imported network looks like this:

```
> net

Boolean network with 5 genes

Involved genes:
Input 1 Input 2 Gene 1 Gene 2 Gene 3
```

```
Transition functions:
Input 1 = 1
Input 2 = Input 2
Gene 1 = (!Gene 1 & !Input 1 & Input 2) | (!Gene 1 & Input 1 & !Input 2)
         | (Gene 1 & !Input 1 & !Input 2) | (Gene 1 & Input 1 & Input 2)
Gene 2 = Gene 1 & Gene 3 & !Input 2
Gene 3 = Gene 1 | Gene 2

Knocked-out and over-expressed genes:
Input 1 = 1
```

We can see that Input 1 is specified as an over-expressed constant gene. Input 2 is modeled as depending only on itself, i.e. it keeps its initial value. Gene 1 is a representation of the XOR function in Disjunctive Normal Form (DNF), using only logical ANDs, logical ORs, and negations. Gene 2 and Gene 3 consist of conjunctions and disjunctions of their inputs respectively. In addition to this textual description, we can visually verify the network by plotting its wiring:

> *plotNetworkWiring(net)*

The resulting plot is shown in Figure 12.

Figure 12: The wiring graph of the imported network specified in BioTapestry.

You can now use the imported network just like any other network in BoolNet.

4.4 Exporting network simulations to Pajek

For further analysis, network simulations can be exported to Pajek, a Windows application that provides visualization and analysis methods for graph structures [3]. For more information on Pajek, please refer to `http://pajek.imfm.si/doku.php`.

The export function writes the state transition graph to a Pajek file (*.net). This requires a synchronous exhaustive attractor search in `BoolNet` to build the full transition table.

To export the mammalian cell cycle network to Pajek, call

```
> data(cellcycle)
> attr <- getAttractors(cellcycle)
> toPajek(attr, file="cellcycle.net")
```

This will export the graph of the state transitions, which is usually sufficient for plotting. If you want to include the state information (i.e., the gene assignment vectors), call

```
> toPajek(attr, file="cellcycle.net", includeLabels=TRUE)
```

Now, start Pajek, load the network with `File | Network | Read`, and check out the tools provided by this application. For example, visualizations can be accessed using the menu item `Draw | Draw`.

Figure 13 shows a plot of the cell cycle network with the Kamada-Kawai layout (Menu entry `Layout | Energy | Kamada-Kawai | Separate Components`).
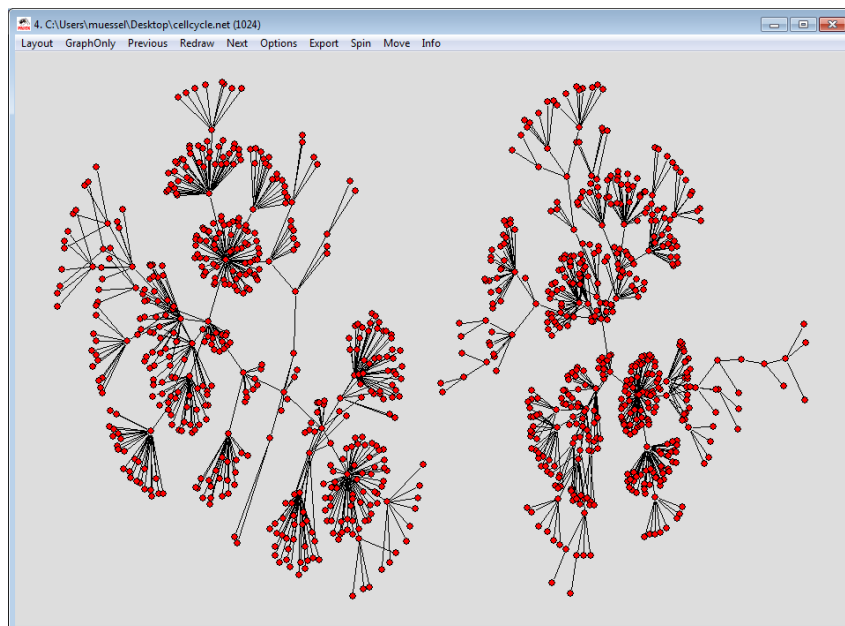

Figure 13: A visualization of the mammalian cell cycle network in Pajek.

References

[1] M. Aldana. Boolean dynamics of networks with scale-free topology. *Physica D*, 185(1):45–66, 2003.

[2] M. Aldana, S. Coppersmith, and L. P. Kadanoff. Boolean dynamics with random coupling. In E. Kaplan, J. E. Marsden, and K. R. Sreenivasan, editors, *Perspectives and Problems in Nonlinear Science*. Springer, 2003.

[3] V. Batagelij and A. Mrvar. Pajek – Program for Large Network Analysis. *Connections*, 21(2):47–57, 1998.

[4] A. Biere. PicoSAT Essentials. *Journal on Satisfiability, Boolean Modeling and Computation*, 4:75–97, 2008.

[5] T. M. Cover and J. A. Thomas. *Information Theory*. Wiley, New York, 1991.

[6] E. Dubrova and M. Teslenko. A SAT-based algorithm for finding attractors in synchronous boolean networks. *IEEE/ACM Transactions on Computational Biology and Bioinformatics*, 8(5):1393–1399, 2011.

[7] A. Fauré, A. Naldi, C. Chaouiya, and D. Thieffry. Dynamical analysis of a generic Boolean model for the control of the mammalian cell cycle. *Bioinformatics*, 22(14):e124–e131, 2006.

[8] S. Kauffman, C. Peterson, B. Samuelsson, and C. Troein. Genetic networks with canalyzing Boolean rules are always stable. *PNAS*, 101(49):17102âĂŞ–17107, 2004.

[9] S. A. Kauffman. Metabolic Stability and Epigensis in Randomly Constructed Genetic Nets. *Journal of Theoretical Biology*, 22(3):437–467, 1969.

[10] S. A. Kauffman. *The Origins of Order: Self-Organization and Selection in Evolution*. Oxford University Press, 1993.

[11] S. Kim, J. Kim, and K.-H. Cho. Inferring gene regulatory networks from temporal expression profiles under time-delay and noise. *Computational Biology and Chemistry*, 31:239–245, 2007.

[12] H. Lähdesmäki, I. Shmulevich, and O. Yli-Harja. On Learning Gene Regulatory Networks Under the Boolean Network Model. *Machine Learning*, 52(1-2):147–167, 2003.

[13] F. Li, T. Long, Q. Ouyang, and C. Tang. The yeast cell-cycle network is robustly designed. *PNAS*, 101:4781–4786, 2004.

[14] S. Liang, S. Fuhrman, and R. Somogyi. REVEAL, a general reverse engineering algorithm for inference of genetic network architectures. *Pacific Symposium on Biocomputing*, 3:18–29, 1998.

[15] W. J. R. Longabaugh, E. H. Davidson, and H. Bolouri. Computational representation of developmental genetic regulatory networks. *Developmental Biology*, 283(1):1–16, 2005.

[16] M. Maucher, B. Kracher, M. Kühl, and H. A. Kestler. Inferring Boolean network structure via correlation. *Bioinformatics*, 27(11):1529–1536, 2011.

[17] I. Shmulevich, E. R. Dougherty, S. Kim, and W. Zhang. Probabilistic Boolean networks: a rule-based uncertainty model for gene-regulatory networks. *Bioinformatics*, 18(2):261–274, 2002.

[18] P. T. Spellman, G. Sherlock, M. Q. Zhang, V. R. Iyer, K. Anders, M. B. Eisen, P. O. Brown, D. Botstein, and B. Futcher. Comprehensive Identification of Cell Cycle-regulated Genes of the Yeast Saccharomyces cerevisiae by Microarray Hybridization. *Molecular Biology of the Cell*, 9(12):3273–3297, 1998.

5 Appendix

5.1 Network file format

This section provides a full language description for the network file format of `BoolNet`. The language is described in Extended Backus-Naur Form (EBNF).

For synchronous, asynchronous and probabilistic Boolean networks, the supported format is as follows:

```
Network           = Header Newline {Rule Newline | Comment Newline};
Header            = "targets" Separator "factors";
Rule              = GeneName Separator BooleanExpression [Separator Probability];
Comment           = "#" String;
BooleanExpression = GeneName
                    | "!" BooleanExpression
                    | "(" BooleanExpression ")"
                    | BooleanExpression " & " BooleanExpression
                    | BooleanExpression " | " BooleanExpression;
                    | "all(" BooleanExpression {"," BooleanExpression} ")"
                    | "any(" BooleanExpression {"," BooleanExpression} ")"
                    | "maj(" BooleanExpression {"," BooleanExpression} ")"
                    | "sumgt(" BooleanExpression {"," BooleanExpression} "," Integer ")"
                    | "sumlt(" BooleanExpression {"," BooleanExpression} "," Integer ")";
GeneName          = ? A gene name from the list of involved genes ?;
Separator         = ",";
Integer           = ? An integer value?;
Probability       = ? A floating-point number ?;
String            = ? Any sequence of characters (except a line break) ?;
Newline           = ? A line break character ?;
```

The extended format for temporal networks includes additional time specifications and temporal predicates is defined as follows:

```
Network              = Header Newline
                        {Function Newline | Comment Newline};
Header               = "targets" Separator "factors";
Function             = GeneName Separator BooleanExpression;
Comment              = "#" String;
BooleanExpression    = GeneName
                        | GeneName TemporalSpecification
                        | BooleanOperator
                        | TemporalOperator
BooleanOperator      =   BooleanExpression
                        | "!" BooleanExpression
                        | "(" BooleanExpression ")"
                        | BooleanExpression " & " BooleanExpression
                        | BooleanExpression " | " BooleanExpression;
TemporalOperator     =   "all" [TemporalIteratorDef]
                               "(" BooleanExpression {"," BooleanExpression} ")"
                        | "any" [TemporalIteratorDef]
                               "(" BooleanExpression {"," BooleanExpression} ")"
                        | "maj" [TemporalIteratorDef]
                               "(" BooleanExpression {"," BooleanExpression} ")"
                        | "sumgt" [TemporalIteratorDef]
                                "(" BooleanExpression {"," BooleanExpression} "," Integer ")"
                        | "sumlt" [TemporalIteratorDef]
                                "(" BooleanExpression {"," BooleanExpression} "," Integer ")"
                        | "timeis" "(" Integer ")"
                        | "timegt" "(" Integer ")"
                        | "timelt" "(" Integer ")";
TemporalIteratorDef  = "[" TemporalIterator "=" Integer ".." Integer "]";
TemporalSpecification = "[" TemporalOperand {"+" TemporalOperand | "-" TemporalOperand} "]";
TemporalOperand      = TemporalIterator | Integer
TemporalIterator     = ? An alphanumeric string ?;
GeneName             = ? A gene name from the list of involved genes ?;
Separator            = ",";
Integer              = ? An integer value?;
String               = ? Any sequence of characters (except a line break) ?;
Newline              = ? A line break character ?;
```